

Olimov Bekhzod (올리모브벙조드)

Kyungpook National University

In [1]:

```
# Principal Component Analysis
from sklearn.datasets import load_iris
from sklearn.decomposition import PCA

# Supervised ML Algorithms
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error, mean_absolute_error

# Data Manipulation
import numpy as np
from numpy.testing import assert_almost_equal
import pandas as pd
import time

# Data Visualization
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
%matplotlib inline
```

The Iris Dataset

This data sets consists of 3 different types of irises' (Setosa, Versicolor, and Virginica) petal and sepal length, stored in a 150x4 numpy.ndarray

The rows being the samples and the columns being: Sepal Length, Sepal Width, Petal Length and Petal Width.

In [2]:

```
# Load iris dataset
iris = load_iris()
X, y = iris.data, iris.target

# Class names
labels = ['Setosa', 'Versicolor', 'Virginica']
# Colors for plotting
colors = ['r', 'g', 'b']
```

In [3]:

```
def plot_by_features(X, y):

    """
    Plots the data by using a pair of features

    Input:

    X -- 2D array dataset with the shape of (number of examples, number of features)
    y -- 1D array containing true labels of the dataset with the shape of (number of examples, )

    Output:

    Scatter plots containing two attributes
    """

    # Seperate each feature by name
    sepal_length = X[:, 0]
```

```

sepal_width = X[:, 1]
petal_length = X[:, 2]
petal_width = X[:, 3]

# The first subplot
plt.figure(figsize=(18, 9))
plt.subplot(2,3,1)
plt.xlabel('Sepal Length')
plt.ylabel('Sepal Width')
plt.xticks(())
plt.yticks(())

for label, color in zip(labels, colors):

    plt.scatter(sepal_length[y==0], sepal_width[y==0], c=colors[0], edgecolor='k')
    plt.scatter(sepal_length[y==1], sepal_width[y==1], c=colors[1], edgecolor='k')
    plt.scatter(sepal_length[y==2], sepal_width[y==2], c=colors[2], edgecolor='k')

plt.legend((labels), fontsize=9, loc='lower right')

# The second subplot
plt.subplot(2,3,2)
plt.xlabel('Sepal Length')
plt.ylabel('Petal Width')
plt.xticks(())
plt.yticks(())

for label, color in zip(labels, colors):

    plt.scatter(sepal_length[y==0], petal_width[y==0], c=colors[0], edgecolor='k')
    plt.scatter(sepal_length[y==1], petal_width[y==1], c=colors[1], edgecolor='k')
    plt.scatter(sepal_length[y==2], petal_width[y==2], c=colors[2], edgecolor='k')

plt.legend((labels), fontsize=9, loc='lower right')

# The third subplot
plt.subplot(2,3,3)
plt.xlabel('Sepal Length')
plt.ylabel('Petal Length')
plt.xticks(())
plt.yticks(())

for label, color in zip(labels, colors):

    plt.scatter(sepal_length[y==0], petal_length[y==0], c=colors[0], edgecolor='k')
    plt.scatter(sepal_length[y==1], petal_length[y==1], c=colors[1], edgecolor='k')
    plt.scatter(sepal_length[y==2], petal_length[y==2], c=colors[2], edgecolor='k')

plt.legend((labels), fontsize=9, loc='lower right')

# The fourth subplot
plt.subplot(2,3,4)
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')
plt.xticks(())
plt.yticks(())

for label, color in zip(labels, colors):

    plt.scatter(petal_length[y==0], petal_width[y==0], c=colors[0], edgecolor='k')
    plt.scatter(petal_length[y==1], petal_width[y==1], c=colors[1], edgecolor='k')
    plt.scatter(petal_length[y==2], petal_width[y==2], c=colors[2], edgecolor='k')

plt.legend((labels), fontsize=9, loc='lower right')

# The fifth subplot
plt.subplot(2,3,5)
plt.xlabel('Petal Length')
plt.ylabel('Sepal Width')
plt.xticks(())
plt.yticks(())

for label, color in zip(labels, colors):

    plt.scatter(petal_length[y==0], sepal_width[y==0], c=colors[0], edgecolor='k')
    plt.scatter(petal_length[y==1], sepal_width[y==1], c=colors[1], edgecolor='k')

```

```

plt.scatter(petal_length[y==2], sepal_width[y==2], c=colors[2], edgecolor='k')

plt.legend((labels), fontsize=9, loc='lower right')

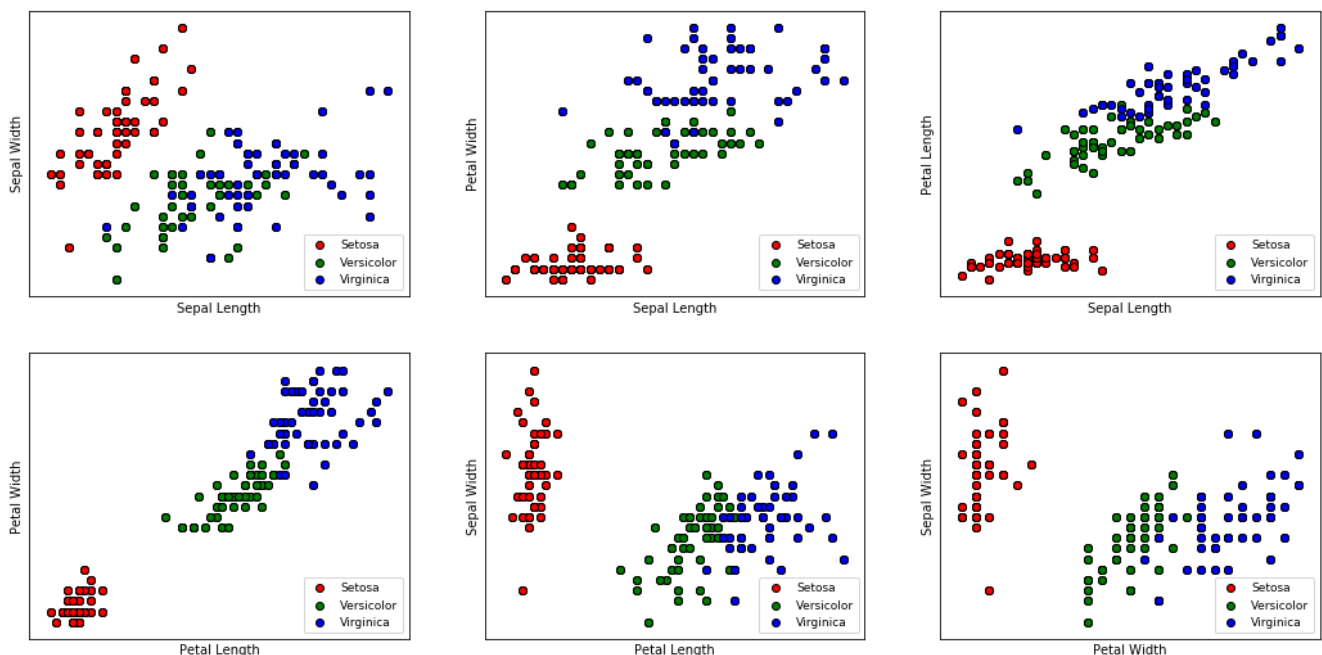
# The sixth subplot
plt.subplot(2,3,6)
plt.xlabel('Petal Width')
plt.ylabel('Sepal Width')
plt.xticks(())
plt.yticks(())

for label, color in zip(labels, colors):

    plt.scatter(petal_width[y==0], sepal_width[y==0], c=colors[0], edgecolor='k')
    plt.scatter(petal_width[y==1], sepal_width[y==1], c=colors[1], edgecolor='k')
    plt.scatter(petal_width[y==2], sepal_width[y==2], c=colors[2], edgecolor='k')

plt.legend((labels), fontsize=9, loc='lower right')
plot_by_features(X, y)

```



Task 1. Perform PCA on Iris dataset

In [4]:

```

def standardization(X):

    """
    Standardizes the data by making it have
    zero mean and unit variance

    Input:

    X -- 2D array dataset with the shape of (number of examples, number of features)

    Output:

    X_standardized -- 2D array standardizaed dataset with the shape of (number of examples, number
    of features)
    """

    mean = np.mean(X, axis=0)
    std = np.std(X, axis=0)
    X_standardized = (X - mean) / std

    print("~~~~~ STANDARDIZATION COMPLETED ~~~~~\n")
    print(f"The first three examples of the original dataset:\n {X[:3]}\n")
    print(f"The first three examples of the standardized dataset:\n {X_standardized[:3]}")

```

```
return X_standardized
```

```
X_standardized = standardization(X)
```

~~~~~ STANDARDIZATION COMPLETED ~~~~~

The first three examples of the original dataset:

```
[[5.1 3.5 1.4 0.2]
 [4.9 3.  1.4 0.2]
 [4.7 3.2 1.3 0.2]]
```

The first three examples of the standardized dataset:

```
[[ -0.90068117  1.01900435 -1.34022653 -1.3154443 ]
 [-1.14301691 -0.13197948 -1.34022653 -1.3154443 ]
 [-1.38535265  0.32841405 -1.39706395 -1.3154443 ]]
```

## Plots

In [5]:

```
def plot_1D(X, y):

    """
    Plots the first two PCA dimensions after applying PCA

    Inputs:

    X -- 2D array dataset with the shape of (number of examples, number of features)
    y -- 1D array containing true labels of the dataset with the shape of (number of examples, )

    Output:

    Graphical illustration of the instances based on the first PC
    """

    # Applying PCA with 1 component
    pca = PCA(n_components=1)
    X_pca = pca.fit_transform(X)
    print(f"Shape of the original data: {X.shape}")
    print(f"Shape of the transformed data: {X_pca.shape}")


    # Plotting
    plt.figure(figsize=(9,5))
    plt.plot(X_pca[y==0], len(X_pca[y==0]) * [1], 'x', c=colors[0])
    plt.plot(X_pca[y==1], len(X_pca[y==1]) * [1], 'x', c=colors[1])
    plt.plot(X_pca[y==2], len(X_pca[y==2]) * [1], 'x', c=colors[2])


    plt.xlabel('The First Principal Component', fontsize=12)
    plt.xlim(-3.5, 3.5)
    plt.legend(labels, loc='lower right', fontsize=11)
    plt.grid()
    plt.yticks(())

plot_1D(X_standardized, y)
```

Shape of the original data: (150, 4)

Shape of the transformed data: (150, 1)





In [6]:

```
def plot_2D(X, y):

    """
    Plots the first two PCA dimensions after applying PCA

    Inputs:

    X -- 2D array standardized dataset with the shape of (number of examples, number of features)
    y -- 1D array containing true labels of the dataset with the shape of (number of examples, )

    Output:

    Graphical illustration of the instances based on the first two PCs
    """
    # Applying PCA with 2 components
    pca = PCA(n_components=2)
    X_pca = pca.fit_transform(X)
    print(f"Shape of the original data: {X.shape}")
    print(f"Shape of the transformed data: {X_pca.shape}")

    # Plotting
    plt.figure(figsize=(9,5))


    plt.scatter(X_pca[:, 0][y == 0], X_pca[:, 1][y == 0],
                c=colors[0], edgecolor='k', s=50)
    plt.scatter(X_pca[:, 0][y == 1], X_pca[:, 1][y == 1],
                c=colors[1], edgecolor='k', s=50)
    plt.scatter(X_pca[:, 0][y == 2], X_pca[:, 1][y == 2],
                c=colors[2], edgecolor='k', s=50)

    plt.xlabel('The First Principal Component', fontsize=12)
    plt.ylabel('The Second Principal Component', fontsize=12)
    plt.xlim(-3.5, 3.5)
    plt.legend(labels, loc='lower right', fontsize=11)
    plt.grid()

plot_2D(X_standardized, y)
```

Shape of the original data: (150, 4)

Shape of the transformed data: (150, 2)



In [7]:

```
def plot_3D(X):

    """
    Plots the first three PCA dimensions after applying PCA
```

Reduce the three three PCA dimensions after applying PCA

Inputs:

X -- 2D array standardized dataset with the shape of (number of examples, number of features)

Output:

Graphical illustration of the instances based on the first three PCs

"""

# Applying PCA with 3 components

X\_pca = PCA(n\_components=3).fit\_transform(X)

print(f"Shape of the original data: {X.shape}")

print(f"Shape of the transformed data: {X\_pca.shape}")

# Plotting

fig = plt.figure(1, figsize=(8, 6))

ax = Axes3D(fig, elev=-150, azim=110)

ax.scatter(X\_pca[:, 0][y==0], X\_pca[:, 1][y==0], X\_pca[:, 2][y==0],  
c=colors[0], edgecolor='k', s=50, label='Setosa')

ax.scatter(X\_pca[:, 0][y==1], X\_pca[:, 1][y==1], X\_pca[:, 2][y==1],  
c=colors[1], edgecolor='k', s=50, label='Versicolor')

ax.scatter(X\_pca[:, 0][y==2], X\_pca[:, 1][y==2], X\_pca[:, 2][y==2],  
c=colors[2], edgecolor='k', s=50, label='Virginica')

ax.set\_xlabel("The first Principal Component")

ax.w\_xaxis.set\_ticklabels([])

ax.set\_ylabel("The second Principal Component")

ax.w\_yaxis.set\_ticklabels([])

ax.set\_zlabel("The third Principal Component")


ax.w\_zaxis.set\_ticklabels([])

ax.legend()

plot\_3D(X\_standardized)

Shape of the original data: (150, 4)

Shape of the transformed data: (150, 3)



## Task 2. Select a proper number of principal components in terms of the proportion of variance

In [8]:

```
def get_optimal_number_of_components_bar(X):
```

"""

*Plots the bar chart to illustrate explained variance ratio with different number of Principal Components*

*Input:*

*X -- 2D array standardized dataset with the shape of (number of examples, number of features)*

*Output:*

*Graphical illustration of the Explained Variance Ratio*

*"""*

*# Applying PCA*

*pca = PCA()*

*X\_pca = pca.fit\_transform(X)*

*# Plotting the explained variances*

*features = range(pca.n\_components\_)*

*var\_ratio = pca.explained\_variance\_ratio\_*

*width = 0.5*

*fig, ax = plt.subplots(nrows=1, ncols=1, figsize=(8,5))*

*rect = ax.bar(features, var\_ratio, color='c', width=width, label='Explained Variances %')*

*# Function to write the values on top of the bars*

*def autolabel(rects):*

*for rect in rects:*

*height = rect.get\_height()*

*ax.annotate('{:.2f}%'.format(height\*100),*

*xy=(rect.get\_x() + rect.get\_width() / 2, height+0.01),*

*ha='center', va='bottom')*

*autolabel(rect)*

*# The graph details*

*plt.xlabel('Components')*

*plt.ylabel('Explained Variances')*

*plt.xlim(-0.5, 3.6)*


*plt.xticks(np.arange(4), labels=['1', '2', '3', '4'])*

*plt.yticks(np.arange(0, 0.85, 0.1), labels=['0', '10', '20', '30', '40', '50', '60', '70', '80'])*

*])*

*plt.grid()*

*get\_optimal\_number\_of\_components\_bar(X\_standardized)*



In [9]:

```
def get_optimal_number_of_components_line_plot(X):
```

*"""*

*Plots the line graph to illustrate lost variance with different number of Principal Components*

*Input:*

*X -- 2D array standardized dataset with the shape of (number of examples, number of features)*

Output:

*Graphical illustration of the "elbow" method*

"""

*# Applying PCA*

pca = PCA().fit(X)

*# Plotting the lost variance*

plt.figure(figsize=(8,5))

plt.plot(np.arange(4), np.cumsum(pca.explained\_variance\_ratio\_), c='b', alpha=0.5,  
marker='^', markersize=7, markerfacecolor='k', linewidth=3, label='Cumulative Sum of R  
etained Variance %')

plt.xlabel('Number of Components')

plt.ylabel('Cumulative Variance %')

plt.grid()

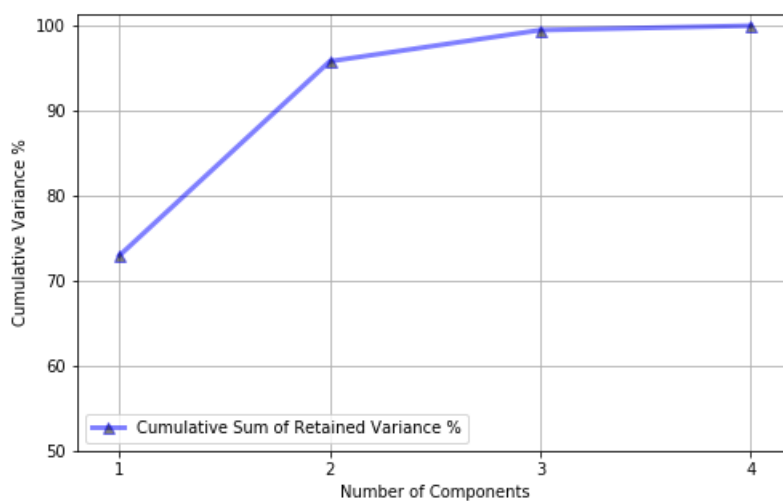
plt.legend()

plt.xticks(np.arange(4), labels=['1', '2', '3', '4'])

plt.xlim(-0.2, 3.2)

plt.yticks(np.arange(0.5, 1.05, 0.1), labels=['50', '60', '70', '80', '90', '100'])

get\_optimal\_number\_of\_components\_line\_plot(X\_standardized)



In [10]:

```
def experiments(X, y, mla='log_reg'):
```

"""

*Applies PCA with different number of components on the data  
and conducts experiments using two types of Machine Learning Algorithms*

*Inputs:*

*X -- 2D array standardized dataset with the shape of (number of examples, number of features)*

*y -- 1D array containing true labels of the dataset with the shape of (number of examples, )*

*mla -- Machine Learning Algorithm to conduct the experiments. Logistic Regression is default*

*Outputs:*

*exp\_df -- a dataframe with the information about the experiments  
with different number of components*

"""

cv\_score\_history = []

time\_history = []

number\_of\_components = []

variance\_history = []

for i in range(1, 5):

tic = time.time()

*# Applying PCA with different number of components*

pca = PCA(n\_components=i)

X\_pca = pca.fit\_transform(X)



```

number_of_components.append(i)
var_ratio = pca.explained_variance_ratio_
variance_history.append("{:.2f}".format(np.cumsum(var_ratio)[-1] * 100))

# Selecting Supervised ML Algorithm
if mla == 'log_reg':
    clf = LogisticRegression(penalty='l2', tol=0.00001,
                             max_iter=300, solver='lbfgs',
                             random_state=2020)

elif mla == 'knn':
    clf = KNeighborsClassifier(n_neighbors=3, weights='uniform',
                              algorithm='auto',
                              leaf_size=20, metric='minkowski')

elif mla == 'svm':
    clf = SVC(C=1, kernel='poly', degree=3,
              gamma='scale', random_state=2020)

elif mla == 'dct':
    clf = DecisionTreeClassifier(criterion='gini', splitter='best',
                                 min_samples_split=2, min_samples_leaf=1,
                                 max_leaf_nodes=20, random_state=2020)

elif mla == 'rfc':
    clf = RandomForestClassifier(n_estimators=50, criterion='gini',
                                 max_depth=5, min_samples_split=3,
                                 min_samples_leaf=1, max_leaf_nodes=10, random_state=2020)

# 5-Fold Cross Validation
cv_scores = cross_val_score(clf, X_pca, y, cv=5)

toc = time.time()

tm = "{:.3f}".format(toc-tic)
time_history.append(tm)

acc = "{:.2f}".format(np.mean(cv_scores) * 100)
cv_score_history.append(acc)

# Creating a Dataframe containing the results
exp_df = pd.DataFrame(data=[number_of_components, variance_history, cv_score_history, time_hist
ory],
                      index=['Num. Components', 'Retained Variance (%)', 'Accuracy Score (%)', 'Time Ela
psed (sec)'],
                      columns=['Exp. 1', 'Exp. 2', 'Exp. 3', 'Exp. 4'])

return exp_df

```

In [11]:

```

df_log = experiments(X_standardized, y)
print("~~~~~ EXPERIMENTS USING Logistic Regression ~~~~~")
print("~~~~~ RESULTS ~~~~~\n")
)
print(df_log)
df_knn = experiments(X_standardized, y, 'knn')
print("\n~~~~~ EXPERIMENTS USING K-Nearest Neighbors ~~~~~")
)
print("~~~~~ RESULTS ~~~~~\n")
)
print(df_knn)
df_svm = experiments(X_standardized, y, 'svm')
print("\n~~~~~ EXPERIMENTS USING Support Vector Classifier
~~~~~")
print("~~~~~ RESULTS
~~~~~\n")
print(df_svm)
df_dct = experiments(X_standardized, y, 'dct')
print("\n~~~~~ EXPERIMENTS USING Decision Tree Classifier
~~~~~")
print("~~~~~ RESULTS
~~~~~\n")
print(df_dct)
df_ran = experiments(X_standardized, y, 'rfc')

```

```
print("\n~~~~~ EXPERIMENTS USING Random Forest Classifier ~~~~~")
print("~~~~~ RESULTS ~~~~~\n")
print(df_ran)
```

```
~~~~~ EXPERIMENTS USING Logistic Regression ~~~~~
~~~~~ RESULTS ~~~~~
```

|                       | Exp. 1 | Exp. 2 | Exp. 3 | Exp. 4 |
|-----------------------|--------|--------|--------|--------|
| Num. Components       | 1      | 2      | 3      | 4      |
| Retained Variance (%) | 72.96  | 95.81  | 99.48  | 100.00 |
| Accuracy Score (%)    | 92.00  | 91.33  | 96.00  | 96.00  |
| Time Elapsed (sec)    | 0.024  | 0.022  | 0.025  | 0.027  |

```
~~~~~ EXPERIMENTS USING K-Nearest Neighbors ~~~~~
~~~~~ RESULTS ~~~~~
```

|                       | Exp. 1 | Exp. 2 | Exp. 3 | Exp. 4 |
|-----------------------|--------|--------|--------|--------|
| Num. Components       | 1      | 2      | 3      | 4      |
| Retained Variance (%) | 72.96  | 95.81  | 99.48  | 100.00 |
| Accuracy Score (%)    | 93.33  | 90.00  | 96.67  | 95.33  |
| Time Elapsed (sec)    | 0.010  | 0.010  | 0.009  | 0.009  |

```
~~~~~ EXPERIMENTS USING Support Vector Classifier ~~~~~
~~~~~ RESULTS ~~~~~
```

|                       | Exp. 1 | Exp. 2 | Exp. 3 | Exp. 4 |
|-----------------------|--------|--------|--------|--------|
| Num. Components       | 1      | 2      | 3      | 4      |
| Retained Variance (%) | 72.96  | 95.81  | 99.48  | 100.00 |
| Accuracy Score (%)    | 89.33  | 90.00  | 92.67  | 92.67  |
| Time Elapsed (sec)    | 0.005  | 0.005  | 0.006  | 0.005  |

```
~~~~~ EXPERIMENTS USING Decision Tree Classifier ~~~~~
~~~~~ RESULTS ~~~~~
```

|                       | Exp. 1 | Exp. 2 | Exp. 3 | Exp. 4 |
|-----------------------|--------|--------|--------|--------|
| Num. Components       | 1      | 2      | 3      | 4      |
| Retained Variance (%) | 72.96  | 95.81  | 99.48  | 100.00 |
| Accuracy Score (%)    | 90.67  | 88.67  | 94.67  | 94.00  |
| Time Elapsed (sec)    | 0.006  | 0.004  | 0.004  | 0.005  |

```
~~~~~ EXPERIMENTS USING Random Forest Classifier ~~~~~
~~~~~ RESULTS ~~~~~
```

|                       | Exp. 1 | Exp. 2 | Exp. 3 | Exp. 4 |
|-----------------------|--------|--------|--------|--------|
| Num. Components       | 1      | 2      | 3      | 4      |
| Retained Variance (%) | 72.96  | 95.81  | 99.48  | 100.00 |
| Accuracy Score (%)    | 90.00  | 89.33  | 92.67  | 93.33  |
| Time Elapsed (sec)    | 0.264  | 0.257  | 0.258  | 0.261  |

### Task 3. Calculate the reconstruction errors from the transformed features

In [12]:

```
def reconstruction_error(X, loss='mse'):

    """
    Computes reconstruction loss after applying PCA with different number of components

    Inputs:

    X -- 2D array standardized dataset with the shape of (number of examples, number of features)
    loss -- type of loss function. 'MSE' is default

    Outputs:

    rl_df -- a dataframe with the information about the reconstruction error
    with different number of components
    bar chart -- a graphical illustration of the reconstruction error
    with different number of components
    """

    m, n = X.shape[0], X.shape[1] # total number of instances and attributes
```

```

loss_history = []

for i in range(1, 5):

    # Applying PCA with different number of components
    pca = PCA(n_components=i)
    X_pca = pca.fit_transform(X) # shape = (150, i)

    # Obtaining the projection onto components in original space
    X_proj = pca.inverse_transform(X_pca) # shape = (150, 4)

    if loss == 'mse':

        # Computing MSE loss
        sum_sq_diffs = np.sum((X_proj - X) ** 2) # float
        mse_loss = (1 / (n * m)) * sum_sq_diffs # float

        # Ensure our MSE loss correctly implemented
        mse_loss_sklearn = mean_squared_error(X, X_proj)
        assert_almost_equal(mse_loss, mse_loss_sklearn)

        loss_history.append("{:.4f}".format(mse_loss))

    elif loss == 'mae':

        # Computing MAE loss
        absolute_difference = np.sum(np.abs(X_proj - X)) # float
        mae_loss = (1 / (n * m)) * absolute_difference # float

        # Ensure our MAE loss correctly implemented
        mae_loss_sklearn = mean_absolute_error(X, X_proj)
        assert_almost_equal(mae_loss, mae_loss_sklearn)

        loss_history.append("{:.4f}".format(mae_loss))

# Transforming into numpy arrays for visualization
loss_history = np.array(loss_history, dtype=np.float32)

# Creating Dataframes with the obtained results
rl_df = pd.DataFrame(loss_history, index=['1 Component', '2 Components', '3 Components', '4 Components'],
                      columns=['Rec. Error'])

##### VISUALIZATION START
#####

# Visualizing the results of MSE loss
fig, ax = plt.subplots(nrows=1, ncols=1, figsize=(8,5))
def autolabel(rects):
    for rect in rects:
        height = rect.get_height()
        ax.annotate(' {:.3f}'.format(height),
                    xy=(rect.get_x() + rect.get_width() / 2, height+0.0009),
                    ha='center', va='bottom')

if loss == 'mse':
    rect = ax.bar(range(4), sorted(loss_history), color='r', alpha=0.5, width=0.5)
    autolabel(rect)
    plt.xlabel('Number of Components', fontsize=12)
    plt.ylabel('Reconstruction MSE Error', fontsize=12)
    plt.xlim(-0.5, 3.6)
    plt.xticks(np.arange(4), labels=['4', '3', '2', '1'])
    plt.yticks(np.arange(0, 0.33, 0.05), labels=['0', '0.05', '0.1', '0.15', '0.2', '0.25', '0.3'])
])
    plt.grid()

elif loss == 'mae':
    rect = ax.bar(range(4), sorted(loss_history), color='g', alpha=0.5, width=0.5)
    autolabel(rect)
    plt.xlabel('Number of Components', fontsize=12)
    plt.ylabel('Reconstruction MAE Error', fontsize=12)
    plt.xlim(-0.5, 3.6)
    plt.xticks(np.arange(4), labels=['4', '3', '2', '1'])
    plt.yticks(np.arange(0, 0.41, 0.05), labels=['0', '0.05', '0.1', '0.15', '0.2', '0.25', '0.3', '0.35', '0.4'])
    plt.grid()

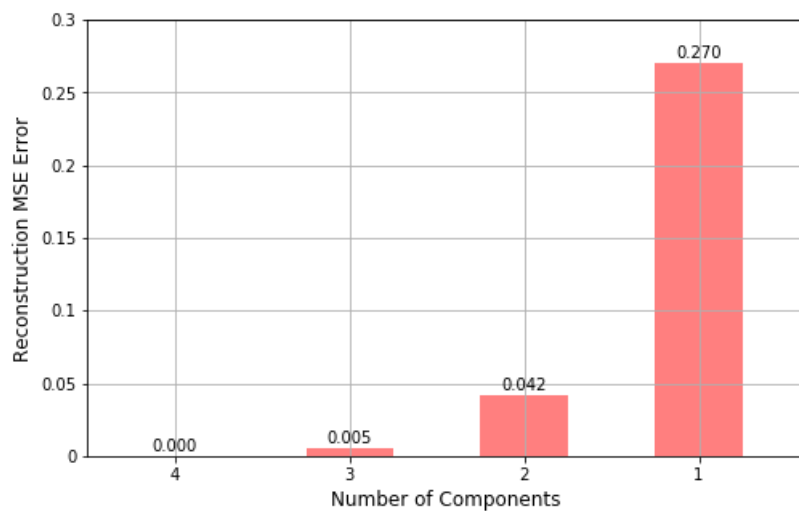
```

```
##### VISUALIZATION END
#####
```

```
return rl_df
```

```
rl_df = reconstruction_error(X_standardized)
print(rl_df)
```

```
Rec. Error
1 Component      0.2704
2 Components      0.0419
3 Components      0.0052
4 Components      0.0000
```



In [13]:

```
rl_df = reconstruction_error(X_standardized, 'mae')
print(rl_df)
```

```
Rec. Error
1 Component      0.3399
2 Components      0.1488
3 Components      0.0467
4 Components      0.0000
```

